

RAG For Open Source LLMs

Presented By

OARC Research Data and Web Platforms Group

Contact Us: webteam@oarc.ucla.edu



Andrew Browning -
Research Data and Web
Platforms Manager



Anthony Doolan - Web
Developer



Hayk Zakaryan -
Software Architect

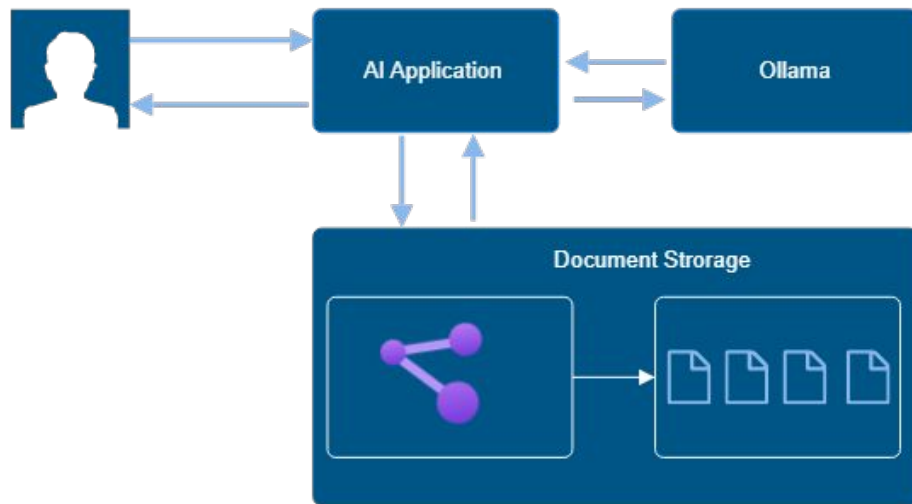
Agenda

1. Introduction to RAG
 2. Application Framework
 3. Ollama Setup & Installation
 4. Prompt Engineering
 5. Sample Python Code
 6. Sample Application Demonstration
-

Introduction to RAG

Retrieval-Augmented Generation (RAG) allows for LLM query retrieval against a stored set of documents restricting LLM inferences to domain specific datasets

Application Framework



- User submits query to AI application
- AI application queries document storage and fetches result
- AI application submits original query and document result to Ollama for inference
- AI application returns result to user

Ollama Setup and Installation

Ollama download page: <https://ollama.com/download>

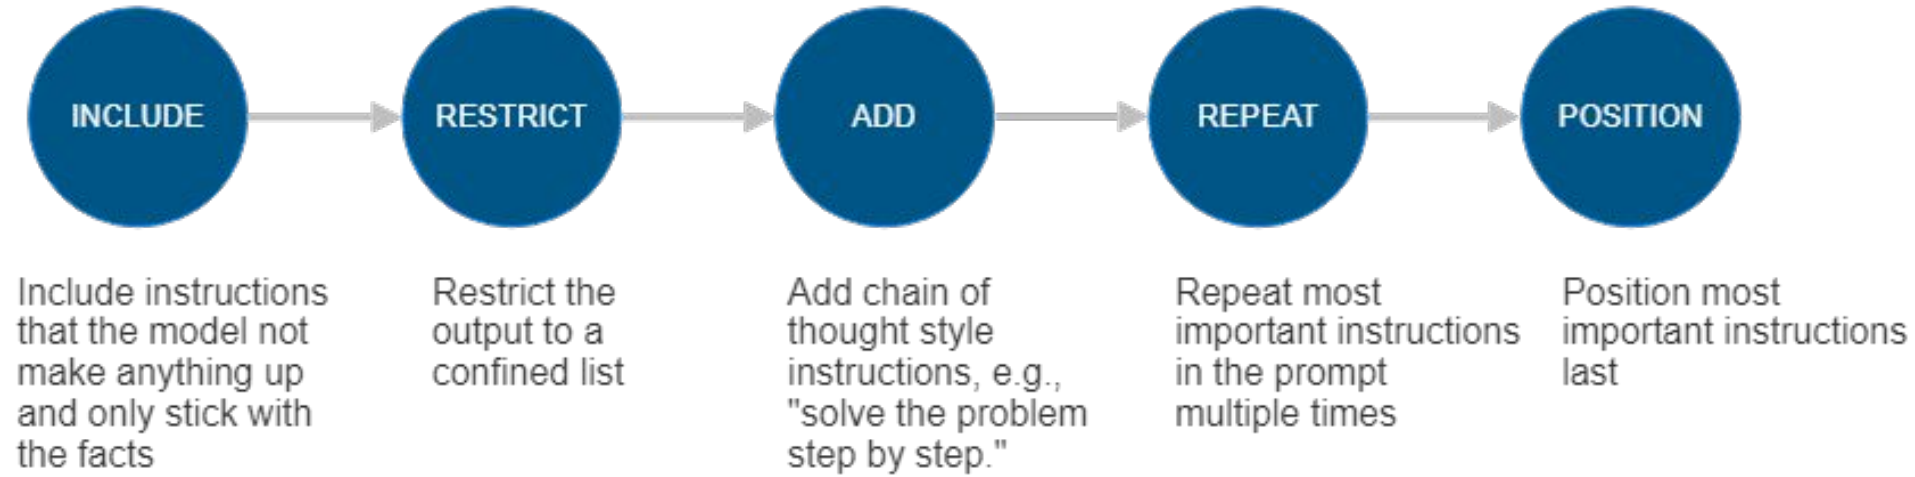
MAC

```
brew install ollama  
ollama serve - start ollama  
ollama pull mistral - download mistral  
ollama list - list available models  
pip install ollama - Install Ollama Python package
```

Linux

```
curl -fsSL https://ollama.com/install.sh | sh  
ollama serve - start ollama  
ollama pull mistral - download mistral  
ollama list - list available models  
pip install ollama - Install Ollama Python package
```

Prompt Engineering



Prompt Engineering

- Prompts guide the LLM to produce the desired output
- Prompt engineering has a huge impact over the LLM output
- Prompt engineering requires creativity, logic, and experimentation
- Prompt engineering requires understanding of capabilities and limitations of the LLM
- Prompt engineering is a tool to control hallucinations
- Prompts are vulnerable to hacking, so meta-prompts must be designed to counteract that

Knowledge Base

- Facebook AI Similarity Search (FAISS)
 - Very easy to setup and experiment with
 - In memory storage
 - Fast
 - Persistent storage is more complicated
- Chroma DB
 - Easy to set up persistent Storage
 - Slower than FAISS for large scale queries

Chunking and Embedding Model

- Ollama and Hugging Face both have embedding models that can be leveraged by langchain.
- Chunking helps to keep the context short and concise.
- Different document types and use cases require different chunking strategies.
- Example Chunking Strategies:
 - Plain Text - Split on character count and new lines
 - CSV - Split on new rows
 - Web Page - Split on headings

Sample Python Code (Prompt)

```
CUSTOM_PROMPT = PromptTemplate(
    template="""
    You are an expert at making desserts. You will only answer questions about making desserts.
    If the user asks you about something else, respond: "I am not able to answer that question."
    If you do not know the answer, say: "I do not know."

    Think about this step by step.

    Answer the following question using only the given context: {context}.
    Question: {question}
    Answer:
    """,
    input_variables=["context", "question"],
)
```

Sample Application Demo

1.) Command line Rag demo -
Uses open source
tools/packages only

2.) UI based Rag demo - Uses
open source tools/packages only



Further Reading - Available Tools and Model Selection

[Ollama](#) - Download page and resources

[Lanchain](#) - Homepage and resources

[Embeddings](#) - Resource to learn more about LLM embeddings

[Python3](#) - Python homepage

[ChromaDB](#) - Chroma homepage and resources

[Marqo](#) - Marqo home page and resources

[Redis](#) - Redis homepage resources